

# MASTERING BLACKJACK WITH REINFORCEMENT LEARNING



# Agenda

1.

## ENVIRONMENT

This section presents the construction of the environment used to model the reinforcement learning agent. It outlines the **observation space**, the set of **available actions**, the **reward structure**, and **specific rules** governing the environment's dynamics..

2.

## BASELINE STRATEGIES

This part presents the implementation of **two baseline strategies**, which serve as **benchmarks** for evaluating the **policy performance** of the reinforcement learning agent.

3.

## FEATURE TRANSFORMATION

This segment covers the development of **two feature transformers** designed to simplify the observation space by **reducing its size** or **capturing nonlinearities** and **potential interactions**.

4.

## AGENT IMPLEMENTATION

This final phase presents the implementation of a reinforcement learning agent using **linear function approximation**. The agent is trained with **model-free algorithms** to learn an effective strategy for playing the game.

# ENVIRONMENT

x x



# Observation Space

32

X

**PLAYER SUM**

**Total value of the player's hand**, ranging from 0 to 31, with face cards counting as 10 and aces as either 1 or 11 depending on usability

11

X

**DEALER UPCARD**

**Value of the dealer's visible face-up card**, ranging from 1 to 11.

(0,1)

X

**USABLE ACE**

Boolean indicating whether the player holds an **ace** that **can be counted as 11** without busting

**Size**

$$\begin{aligned}32 \times 11 \times 2 \\= 704\end{aligned}$$

C U S T O M I Z E D

# Observation Space

32

×

## PLAYER SUM

Total value of the player's hand, ranging from 0 to 31, with face cards counting as 10 and aces as either 1 or 11 depending on usability

11

×

## DEALER UPCARD

Value of the dealer's visible face-up card, ranging from 1 to 11.

(0,1)

×

## USABLE ACE

Boolean indicating whether the player holds an ace that can be counted as 11 without busting

11

×

## PEEKED CARD


Value of the peeked card, ranging from 1 (ace) to 10.

## Size

$$32 \times 11 \times 2 \times 11 \\ = 7744 !$$


G Y M P A C K A G E

# Action Space



**HIT**

The player chooses to **draw another card** from the deck.




**STAND**

The player chooses to **stop drawing cards** and **end their turn**.


C U S T O M I Z E D

# Action Space




**HIT**

The player chooses to **draw another card** from the deck.



**STAND**

The player chooses to **stop drawing cards** and **end their turn**.



**PEEK**

The player **attempts** to **look at the next card** in the deck; if **successful**, it provides **valuable information** for deciding whether to hit or stand, but if **caught**, it results in a **penalty**.

# Rules

## SHUFFLING RULES

- The deck is reshuffled at the **start of the game**.
- As the game progresses, drawn cards are **not returned** to the deck and **cannot appear again**.
- Once **fewer than 30%** of the original cards remain, the **deck is reshuffled**.




## PEEKING RISK

- The player is assumed to be an **expert gambler**, **successfully peeking** at cards **85% of the time**.
- However, each time the player is **caught peeking**, the dealer becomes more watchful, **reducing** the player's future **success rate by 40%** per **failed attempt**.


S C E N A R I O 1

# Reward Structure




**HIT**

The player receives a reward of **0** if the action **does not result in a bust**, and **-1** if it does, indicating the **game is lost**.



**STAND**


The player receives **+1** for **winning** the game, **-1** for **losing**, **0** for **draw**




**PEEK**

The player receives **0** for **successfully peeking**, and incurs a **penalty** of **-0.5** if the peek is **unsuccessful**.


# Reward Structure

**HIT**

The player receives a reward of **0** if the action **does not result in a bust**, and **-1** if it does, indicating the **game is lost**.


**STAND**

The player receives **+1** for **winning** the game, **-1** for **losing**, **0** for **draw**


**PEEK**

The player receives **0** for **successfully peeking**, and incurs a **penalty** of **-2** if the peek is **unsuccessful**.


# Reward Structure

**HIT**

The player receives a reward of **0** if the action **does not result in a bust**, and **-1** if it does, indicating the **game is lost**.

**STAND**

The player receives **+1** for **winning** the game, **-1** for **losing**, **0** for **draw**

**PEEK**

The player receives **0** for **successfully peeking**, and incurs a **penalty** of **-1** if the peek is **unsuccessful**.

# BASELINE STRATEGIES

X X

# Baseline Strategies

## RANDOM

- The player chooses **randomly** between the actions: **hit**, **stand**, and **peek**.



## BASIC


- The player follows a basic strategy: **80% of the time**, they use a naïve approach, **hitting** when their **card total** is **17 or less**, and **standing** if it is **18 or higher**.
- In the **remaining 20%**, the player uses **peeking**. If the peek is **successful**, they **hit** only if the **sum of their cards and the peeked card** is **21 or less**; otherwise, they **stand**.

# FEATURE TRANSFORMATION

× ×




# Feature Transformation



## THRESHOLD

- Converts game state into **categorical levels** (**Low, Medium, High**) mapped to **float values**.
- Focuses on **strategic abstractions** like usable ace and dealer danger level.
- Embeds the state only for the selected action using **block-wise encoding**.



## POLYNOMIAL

- **Normalizes** raw state values and applies a **degree-2 polynomial expansion**.
- Encodes action using **one-hot vectors** and **captures all pairwise feature interactions**.
- Focuses on **non-linear** modeling of complex game patterns.

# Threshold Transformation

EXAMPLE

PLAYER SUM:  14

DEALER UPCARD: 

USABLE ACE: 

PEEKED CARD: 

ACTION: 

# Threshold Transformation

EXAMPLE

PLAYER SUM:  14

DEALER UPCARD: 



USABLE ACE: 

PEEKED CARD: 

ACTION: 

0.	0.	0.	0.	Hit
0.	0.	0.	0.	Stand
<b>0.5</b>	<b>1.</b>	<b>0.</b>	<b>-0.5</b>	Peek
Medium	High	Usable Ace	No Card	

# Polynomial Transformation

$$\tilde{s}_i = \frac{s_i - \min_i}{\max_i - \min_i}$$

## NORMALIZATION

Each raw state value  $s_i$  is **linearly scaled** to the interval  $[0,1]$ , where  $\min_i$  and  $\max_i$  are the predefined bounds for that feature.

$$e(a)_k = \begin{cases} 1, & k = a, \\ 0, & k \neq a. \end{cases}$$

## ACTION ENCODING

The integer **action**  $a \in \{0,1,2\}$  is converted to a **one-hot vector**  $e(a) \in \{0,1\}^3$

$$\phi(\mathbf{x}) = \{x_i x_j \mid 0 \leq i \leq j < d\}$$

## FEATURE EXPANSION

A **bias term 1**, the **normalized state**  $\tilde{\mathbf{s}}$ , and the **one-hot action**  $e(a)$  are concatenated into a **vector**  $\mathbf{x}$ . Then every **pairwise** (including squared) **product** is computed, where  $d$  is the length of  $x$ .

# AGENT IMPLEMENTATION

× ×



# Agent Implementation

$$Q(s, a) = w \cdot \phi(s, a)$$

## Q-VALUES FOR $(s, a)$

The agent begins by estimating how good it is to take a particular action in a given state by projecting the state-action features onto its weight vector. In practice, the feature transformer produces a **vector**  $\phi(s, a)$  and the **inner product** with  $w$  yields the **current value estimate**



$$w \leftarrow w + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)] \phi(s, a)$$

## Q-LEARNING UPDATE

After executing an action and **observing reward**  $r$  and **successor state**  $\tilde{s}$ , the agent computes the target as the **immediate reward plus the discounted maximal estimated future value**. The difference between this target and the prior estimate serves as an **error signal**, which when scaled by the **learning rate**  $\alpha$  and the **feature vector**  $\phi(s, a)$  produces the **weight adjustment**

1

2

$$w \leftarrow w + \alpha [r + \gamma Q(s', a') - Q(s, a)] \phi(s, a)$$

## SARSA UPDATE



The update similarly hinges on the temporal-difference error, but uses the value of the **action**  $\tilde{a}$  actually selected in  $s'$  rather than the best possible. By doing so, the **agent's exploration choices directly shape the learning target**

# Training Evaluation

## Q-Learning      SARSA



Peak Reward = -0.5

Feature Transformation	Average Return	Win Rate	Average Return	Win Rate
Threshold	-0.1263	49.21%	-0.1257	49.20%
Polynomial	-0.1329	49.49%	-0.1401	49.08%





Peak Reward = -1

Feature Transformation	Average Return	Win Rate	Average Return	Win Rate
Threshold	-0.1417	48.37%	-0.14	48.45%
Polynomial	-0.1469	48.47%	-0.1542	48.12%



Peak Reward = -2

Feature Transformation	Average Return	Win Rate	Average Return	Win Rate
Threshold	-0.1907	40.23%	-0.1891	40.33%
Polynomial	-0.2108	39.39%	-0.2119	39.33%



# Training Evaluation

Q-Learning

SARSA

Peak Reward = -0.5


Feature Transformation	Average Return	Win Rate	Average Return	Win Rate
Threshold	-0.1263	49.21%	<b>-0.1257</b>	<b>49.20%</b>
Polynomial	-0.1329	49.49%	-0.1401	49.08%

Peak Reward = -1



Feature Transformation	Average Return	Win Rate	Average Return	Win Rate
Threshold	-0.1417	48.37%	-0.14	48.45%
Polynomial	-0.1469	48.47%	-0.1542	48.12%

Peak Reward = -2



Feature Transformation	Average Return	Win Rate	Average Return	Win Rate
Threshold	-0.1907	40.23%	-0.1891	40.33%
Polynomial	-0.2108	39.39%	-0.2119	39.33%



# Final Comparison



# Final Comparison





THANK YOU